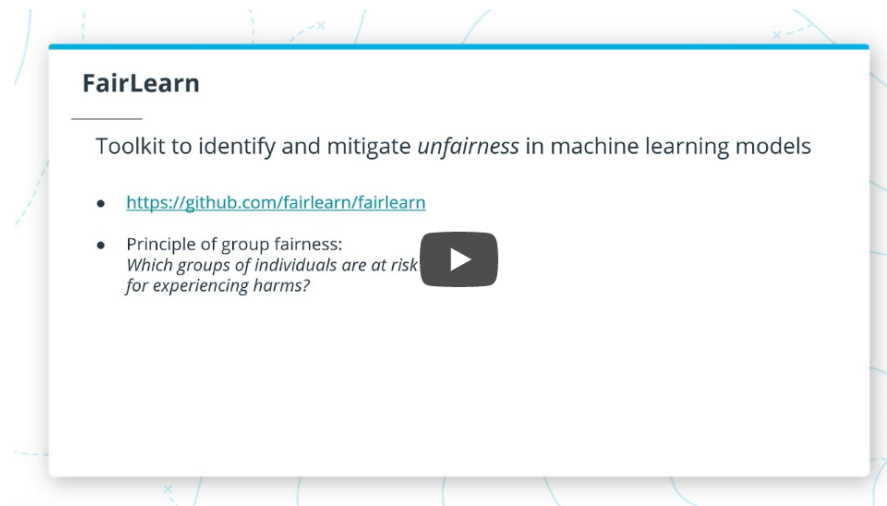


Model Fairness

A presentation slide titled 'FairLearn' with a blue header. The slide describes it as a 'Toolkit to identify and mitigate unfairness in machine learning models'. It includes a link to the GitHub repository and a bullet point about the principle of group fairness, which asks 'Which groups of individuals are at risk for experiencing harms?'. A play button icon is next to the bullet point.

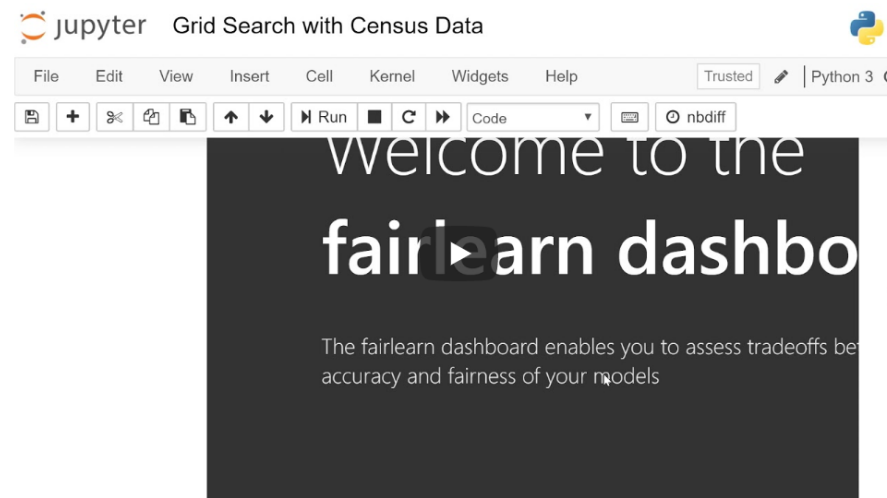
FairLearn

Toolkit to identify and mitigate *unfairness* in machine learning models

- <https://github.com/fairlearn/fairlearn>
- Principle of group fairness:
Which groups of individuals are at risk for experiencing harms?

In this next video, we'll walk through an example of how you could use some of the [Fairlearn](#) capabilities when applying model fairness to your work. This demo is just intended as an example, not a lab exercise, although you are certainly welcome to check out the FairLearn repository for yourself, including the notebooks we use in the demo, [here](#).

Fairlearn Notebook Demo

A screenshot of a Jupyter Notebook interface. The title bar says 'jupyter Grid Search with Census Data'. The menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. The toolbar shows icons for saving, adding, deleting, and running cells, along with a 'Run' button and a 'Code' dropdown. The main content area displays a dark-themed slide titled 'welcome to the fairlearn dashboard'. The slide text says: 'The fairlearn dashboard enables you to assess tradeoffs between accuracy and fairness of your models'.

jupyter Grid Search with Census Data

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

Run Code nbdiff

welcome to the
fairlearn dashboard

The fairlearn dashboard enables you to assess tradeoffs between accuracy and fairness of your models